

**G. H. Rasoni Institute of Engineering & Management, Pune**

**DEPARTMENT: ARTIFICIAL INTELLIGENCE**

**ACADEMIC YEAR: 2022-23**

**CLASS: T.Y. BTech SEMESTER: VI**

**Subject Name: BIG DATA COMPUTING**

**LAB MANUAL**

<b>Sr. No.</b>	<b>Name of the Experiments</b>
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	

**Assignment 1**

**Title:** To install Hadoop framework, configure it and setup a single node cluster. Use web based tools to monitors your Hadoop setup.

**SOFTWARE REQUIREMENTS:**

1.Ubuntu

## 2. jdk 8

### **THEORY:**

Installation of Hadoop:

Hadoop software can be installed in three modes of operation

- **Stand Alone Mode:** Hadoop is a distributed software and is designed to run on a commodity of machines. However, we can install it on a single node in stand-alone mode. In this mode, Hadoop software runs as a single monolithic java process. This mode is extremely useful for debugging purpose. You can first test run your Map-Reduce application in this mode on small data, before actually executing it on cluster with big data.
- **Pseudo Distributed Mode:** In this mode also, Hadoop software is installed on a Single Node. Various daemons of Hadoop will run on the same machine as separate java processes. Hence all the daemons namely NameNode, DataNode, SecondaryNameNode, JobTracker, TaskTracker run on single machine.
- **Fully Distributed Mode:** In Fully Distributed Mode, the daemons NameNode, JobTracker, SecondaryNameNode (Optional and can be run on a separate node) run on the Master Node. The daemons DataNode and TaskTracker run on the Slave Node.

.....Single Node Installation.....

```
sudo apt-get update
```

```
ssh-keygen
```

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

-----\*\*downloading and installing java (Hadoop requires a working Java 1.6+)\*\*-----

```
sudo apt-get install openjdk-8-jdk -y
```

```
-----**downloading hadoop**-----
```

```
wget https://archive.apache.org/dist/hadoop/common/hadoop-1.2.1/hadoop-1.2.1.tar.gz
```

```
tar -xzvf hadoop-1.2.1.tar.gz
```

```
sudo mv hadoop-1.2.1 /usr/local/hadoop
```

```
-----**configuring bashrc linux env setup**-----
```

```
#Configuring bashrc linux env
```

```
nano ~/.bashrc
```

```
export HADOOP_PREFIX=/usr/local/hadoop/
```

```
export PATH=$PATH:$HADOOP_PREFIX/bin
```

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

```
export PATH=$PATH:$JAVA_HOME
```

```
#Setting hadoop env
```

```
nano /usr/local/hadoop/conf/hadoop-env.sh
```

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

```
export HADOOP_OPTS=-Djava.net.preferIPv4Stack=true
```

```
#Configuring core-site.xml
```

```
nano /usr/local/hadoop/conf/core-site.xml
```

```
<property>
```

```
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
```

```
<property>
<name>hadoop.tmp.dir</name>
<value>/usr/local/hadoop/tmp</value>
</property>
```

#Configuring hdfs-site.xml

nano /usr/local/hadoop/conf/hdfs-site.xml

```
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
```

#Configuring mapred-site.xml

nano /usr/local/hadoop/conf/mapred-site.xml

```
<property>
<name>mapred.job.tracker</name>
<value>hdfs://localhost:9001</value>
</property>
```

-----\*\*making tmp dir\*\*-----

mkdir /usr/local/hadoop/tmp

-----\*\*exec bash\*\*-----

exec bash

-----\*\*formatting hadoop namenode\*\*-----

hadoop namenode -format

-----\*\*starting hadoop services(daemons)\*\*-----

start-dfs.sh

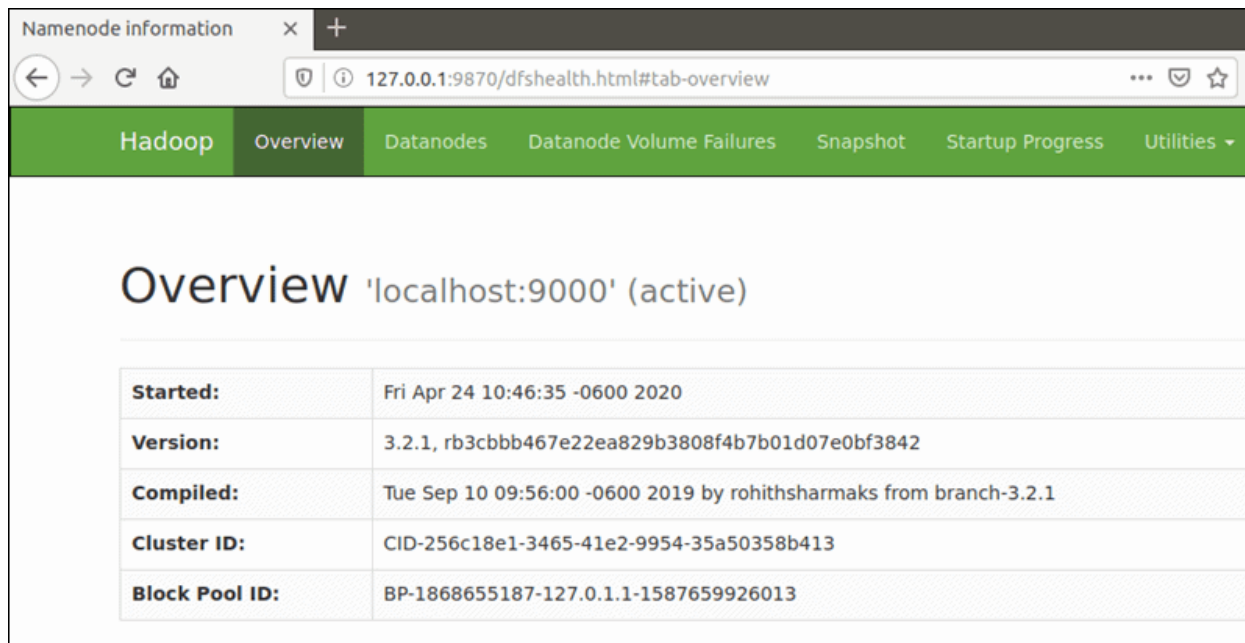
start-mapred.sh

jps

## Access Hadoop UI from Browser

<http://localhost:9870>

The NameNode user interface provides a comprehensive overview of the entire cluster.



The screenshot shows a web browser window with the title 'Namenode information'. The address bar displays '127.0.0.1:9870/dfshealth.html#tab-overview'. The page has a green navigation bar with tabs: 'Hadoop', 'Overview' (selected), 'Datanodes', 'Datanode Volume Failures', 'Snapshot', 'Startup Progress', and 'Utilities'. The main content area is titled 'Overview 'localhost:9000' (active)'. Below the title is a table with the following information:

<b>Started:</b>	Fri Apr 24 10:46:35 -0600 2020
<b>Version:</b>	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
<b>Compiled:</b>	Tue Sep 10 09:56:00 -0600 2019 by rohithsharmaks from branch-3.2.1
<b>Cluster ID:</b>	CID-256c18e1-3465-41e2-9954-35a50358b413
<b>Block Pool ID:</b>	BP-1868655187-127.0.1.1-1587659926013

The default port **9864** is used to access individual DataNodes directly from your browser:

http://localhost:9864

Namenode InformationDataNode InformationAll Applications

127.0.1.1:9864/datanode.html

HadoopOverviewUtilities

DataNode on pnap-VirtualBox:9866

Cluster ID:CID-256c18e1-3465-41e2-9954-35a50358b413

Version:3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842

Block Pools

Namenode Address	Block Pool ID	Actor State	Last Heartbeat	Last Block Report	Last Block Report Size (Max Size)
localhost:9000	BP-1868655187-127.0.1.1-1587659926013	RUNNING	0s	30 minutes	0 B (64 MB)


The YARN Resource Manager is accessible on port **8088**:

http://localhost:8088

The Resource Manager is an invaluable tool that allows you to monitor all running processes in your Hadoop cluster.

Namenode InformationDataNode InformationAll Applications

127.0.1.1:8088/cluster



Cluster

- About Nodes
- Node Labels
- Applications
  - NEW
  - SAVING
  - SUBMITTED
  - ACCEPTED
  - RUNNING
  - FINISHED
  - FAILED
  - KILLED
- Scheduler

Tools

All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved
0	0	0	0	0	0 B	8 GB	0 B

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes
1	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[memory-mb (unit=Mb), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>

Show 20 entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCoers	Allocated Memory MB	Reserved CPU VCoers
No data available in table														

Showing 0 to 0 of 0 entries

## Conclusion :

### Assignment 2

**Title:** To implement file management task in Hadoop HDFS like adding, retrieving, and deleting files.

#### **SOFTWARE REQUIREMENTS:**

1. Ubuntu
2. jdk8

#### **THEORY:**

Implement the following file management tasks in Hadoop:

- i. Adding files and directories

- ii. Retrieving files
- iii. Deleting files

### **Adding Files and Directories to HDFS**

```
hadoop fs -mkdir /user/chuck
```

```
hadoop fs -put
```

```
hadoop fs -put example.txt /user/chuck
```

### **Retrieving Files from HDFS**

The Hadoop command get copies files from HDFS back to the local filesystem. To retrieve example.txt, we can run the following command:

```
hadoop fs -cat
```

```
example.txt
```

### **Deleting files from HDFS**

```
hadoop fs -rm example.txt
```

- Command for creating a directory in hdfs is “hdfs dfs –mkdir /lendicse”.
- Adding directory is done through the command “hdfs dfs –put lendi\_english /”.



## 4.1 INPUT/OUTPUT:

Browsing HDFS - Mozilla Firefox

Browsing HDFS

localhost:50070/explorer.html#/

Hadoop Overview Datanodes Snapshot Startup Progress Utilities -

### Browse Directory

/ Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	lendi	supergroup	0 B	Wed 17 Aug 2016 02:44:00 AM EDT	0	0 B	<a href="#">lendi_english</a>
drwxr-xr-x	lendi	supergroup	0 B	Wed 17 Aug 2016 02:17:48 AM EDT	0	0 B	<a href="#">sadhana</a>
drwxr-xr-x	lendi	supergroup	0 B	Sat 13 Aug 2016 01:31:42 AM EDT	0	0 B	<a href="#">shakes</a>
drwxr-xr-x	lendi	supergroup	0 B	Sat 13 Aug 2016 01:35:59 AM EDT	0	0 B	<a href="#">shakes1</a>
drwx-----	lendi	supergroup	0 B	Sat 13 Aug 2016 01:19:03 AM EDT	0	0 B	<a href="#">tmp</a>

## Assignment 3

**Title:** To implement a word count application using the MapReduce API.

### SOFTWARE REQUIREMENTS:

1. Ubuntu
2. jdk 8

### THEORY:

WordCount is a simple program which counts the number of occurrences of each word in a given text input data set. WordCount fits very well with the MapReduce programming model making it a great example to understand the Hadoop Map/Reduce programming style. Our implementation consists of three main parts:

1. Mapper
2. Reducer
3. Driver

Step-1. Write a Mapper

A Mapper overrides the `map()` function from the Class `"org.apache.hadoop.mapreduce.Mapper"` which provides `<key, value>` pairs as the input. A Mapper implementation may output

`<key,value>` pairs using the provided Context .

Input value of the WordCount Map task will be a line of text from the input data file and the key would be the line number `<line_number, line_of_text>` . Map task outputs `<word, one>` for each word in the line of text.

### ***Pseudo-code***

```
void Map (key, value){  
    for each word x in  
        value:
```

```
output.collect(x,1);
```

## Step-2. Write a Reducer

A Reducer collects the intermediate <key,value> output from multiple map tasks and assemble a single result.

Here, the WordCount program will sum up the occurrence of each word to pairs as

<word, occurrence>.

Pseudo-code

```
void Reduce (keyword, <list of value>){ for  
  
    each x in <list of value>:  
  
        sum+=x;  
  
    final_output.collect(keyword,  
  
    sum);  
  
}
```

## 5.2 INPUT/OUTPUT:

```
lendi@ubuntu: ~/Desktop
16/08/17 01:17:45 INFO impl.YarnClientImpl: Submitted application application_1471410736896_0001
16/08/17 01:17:45 INFO mapreduce.Job: The url to track the job: http://ubuntu.ubuntu-domain:8088/proxy/application_1471410736896_0001/
16/08/17 01:17:45 INFO mapreduce.Job: Running job: job_1471410736896_0001
16/08/17 01:17:52 INFO mapreduce.Job: Job job_1471410736896_0001 running in uber mode : false
16/08/17 01:17:52 INFO mapreduce.Job:  map 0% reduce 0%
16/08/17 01:17:59 INFO mapreduce.Job:  map 100% reduce 0%
16/08/17 01:18:06 INFO mapreduce.Job:  map 100% reduce 100%
16/08/17 01:18:06 INFO mapreduce.Job: Job job_1471410736896_0001 completed successfully
16/08/17 01:18:06 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=3772644
        FILE: Number of bytes written=7775215
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=174718
        HDFS: Number of bytes written=510970
        HDFS: Number of read operations=6
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
```

## Assignment 7

### PROGRAM LOGIC:

#### STEPS FOR INSTALLING APACHE PIG

- 1) Extract the pig-0.15.0.tar.gz and move to home directory
- 2) Set the environment of PIG in bashrc file.
- 3) Pig can run in two modes Local Mode and Hadoop Mode Pig -x local and pig
- 4) Grunt

Shell Grunt >

5) LOADING Data into Grunt Shell

DATA = LOAD <CLASSPATH> USING PigStorage(DELIMITER) as (ATTRIBUTE :  
DataType1, ATTRIBUTE : DataType2.....)

6) Describe

Data Describe

DATA;

7) DUMP

Data Dump

DATA;

## 8.1 INPUT/OUTPUT:

Input as Website Click Count Data

```
lendl@ubuntu: ~  
grunt> ad1 = load '/home/lendl/Desktop/static_data/ad_data/ad_data1.txt' using PigStorage('\t') as (item:chararray,campaignId:chararray,date:chararray,time:chararray,display_site:chararray,was_clicked:int,cpc:int,country:chararray,placement:chararray);  
2016-10-14 02:35:32,441 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2016-10-14 02:35:32,441 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS  
grunt> describe ad1;  
ad1: {item: chararray,campaignId: chararray,date: chararray,time: chararray,display_site: chararray,was_clicked: int,cpc: int,country: chararray,placement: chararray}  
grunt> ad2 = load '/home/lendl/Desktop/static_data/ad_data/ad_data2.txt' using PigStorage(',') as (campaignId:chararray,date:chararray,time:chararray,display_site:chararray,placement:chararray,was_clicked:int,cpc:int,item:chararray);  
2016-10-14 02:36:08,732 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2016-10-14 02:36:08,732 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS  
grunt> describe ad2;  
ad2: {campaignId: chararray,date: chararray,time: chararray,display_site: chararray,placement: chararray,was_clicked: int,cpc: int,item: chararray}  
grunt>
```

## PROGRAM LOGIC:

### FILTER Data

FDATA = FILTER DATA by ATTRIBUTE = VALUE;

### GROUP Data

GDATA = GROUP DATA by ATTRIBUTE;

### Iterating Data

FOR\_DATA = FOREACH DATA GENERATE GROUP AS GROUP\_FUN,  
ATTRIBUTE = <VALUE>

### Sorting Data

SORT\_DATA = ORDER DATA BY ATTRIBUTE WITH CONDITION;

### LIMIT Data

LIMIT\_DATA = LIMIT DATA COUNT;

### JOIN Data

JOIN DATA1 BY (ATTRIBUTE1,ATTRIBUTE2....) , DATA2 BY  
(ATTRIBUTE3,ATTRIBUTE....N)

## 9.1 INPUT / OUTPUT :



```
lendl@ubuntu: ~  
grunt> join_data = join ad1 by (campaignId,display_site,cpc),ad2 by (campaignId,  
display_site,cpc);  
grunt> describe join_data;  
join_data: {ad1::item: chararray,ad1::campaignId: chararray,ad1::date: chararray  
,ad1::time: chararray,ad1::display_site: chararray,ad1::was_clicked: int,ad1::cp  
c: int,ad1::country: chararray,ad1::placement: chararray,ad2::campaignId: charar  
ray,ad2::date: chararray,ad2::time: chararray,ad2::display_site: chararray,ad2::  
placement: chararray,ad2::was_clicked: int,ad2::cpc: int,ad2::item: chararray}  
grunt> 
```

## EXPERIMENT NO.4



## Creating the HDFS tables and loading them in Hive

HIVE tables provide us the schema to store data in various formats (like CSV). Hive provides multiple ways to add data to the tables. We can use DML(Data Manipulation Language) queries in Hive to import or add data to the table. One can also directly put the table into the hive with [HDFS commands](#). In case we have data in Relational Databases like MySQL, ORACLE, IBM DB2, etc. then we can use ***Sqoop*** to efficiently transfer PetaBytes of data between Hadoop and Hive. In this particular tutorial, we will be using Hive DML queries to Load or INSERT data to the Hive table.

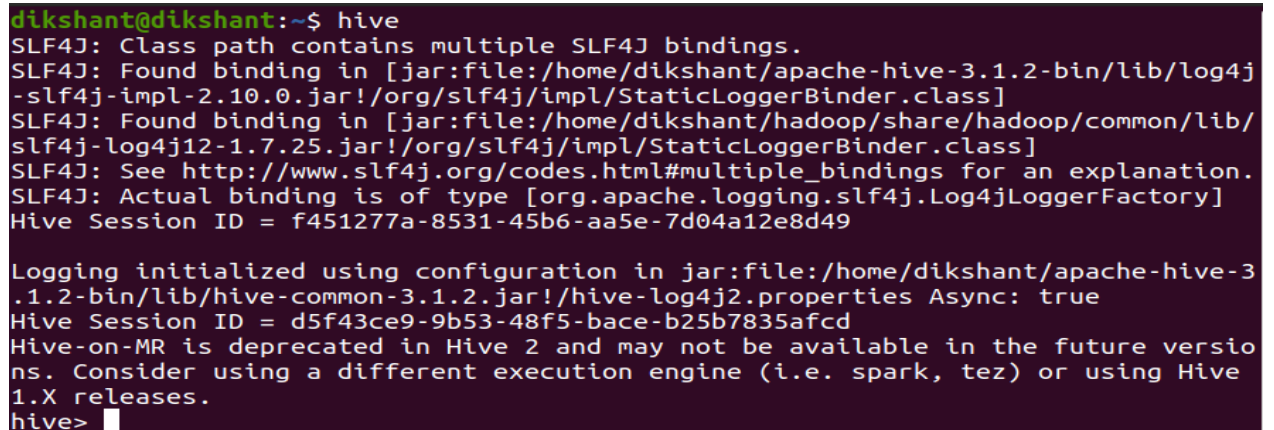
To perform the below operation make sure your hive is running. Below are the steps to launch a hive on your local system.

### Step 1: Start all your Hadoop Daemon

```
start-dfs.sh          # this will start namenode, datanode and secondary namenode
start-yarn.sh         # this will start node manager and resource manager
jps                  # To check running daemons
```

### Step 2: Launch hive from terminal

hive

A terminal window with a dark background showing the output of the 'hive' command. The prompt is 'dikshant@dikshant:~\$'. The output includes SLF4J logging messages about multiple bindings, the Hive Session ID 'f451277a-8531-45b6-aa5e-7d04a12e8d49', and a deprecation warning for 'Hive-on-MR'. The prompt 'hive>' is visible at the bottom.

```
dikshant@dikshant:~$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/dikshant/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/dikshant/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = f451277a-8531-45b6-aa5e-7d04a12e8d49

Logging initialized using configuration in jar:file:/home/dikshant/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = d5f43ce9-9b53-48f5-bace-b25b7835afcd
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> █
```

In hive with DML statements, we can add data to the Hive table in 2 different ways.

- Using INSERT Command
- Load Data Statement

### 1. Using INSERT Command

**Syntax:**

INSERT INTO TABLE <table\_name> VALUES (<add values as per column entity>);

**Example:**

To insert data into the table let's create a table with the name *student* (By default hive uses its *default* database to store hive tables).

**Command:**

```
CREATE TABLE IF NOT EXISTS student(
```

```
Student_Name STRING,
```

```
Student_Rollno INT,
```

```
Student_Marks FLOAT)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ',';
```

```
hive> CREATE TABLE IF NOT EXISTS student(
> Student_Name STRING,
> Student_Rollno INT,
> Student_Marks FLOAT)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ',';
OK
Time taken: 2.086 seconds
hive> show tables in default;
OK
student
Time taken: 0.268 seconds, Fetched: 1 row(s)
hive> █
```

We have successfully created the *student* table in the Hive *default* database with the attribute *Student\_Name*, *Student\_Rollno*, and *Student\_Marks* respectively.

Now, let's insert data into this table with an INSERT query.

**INSERT Query:**

```
INSERT INTO TABLE student VALUES ('Dikshant',1,'95'),('Akshat', 2 , '96'),('Dhruv',3,'90');
```

```
hive> INSERT INTO TABLE student VALUES ('Dikshant',1,'95'),('Akshat', 2 , '96'),
('Dhruv',3,'90');
Query ID = dikshant_20201106121659_f5dfa694-f552-4b7a-a64b-4f3804213ab8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
```

We can check the data of the *student* table with the help of the below command.

```
SELECT * FROM student;
```

```
hive> SELECT * FROM student;
OK
Dikshant      1      95.0
Akshat  2      96.0
Dhruv   3      90.0
Time taken: 0.162 seconds, Fetched: 3 row(s)
hive> █
```

## 2. Load Data Statement

Hive provides us the functionality to load pre-created table entities either from our local file system or from HDFS. The *LOAD DATA* statement is used to load data into the hive table.

### Syntax:

```
LOAD DATA [LOCAL] INPATH '<The table data location>' [OVERWRITE] INTO TABLE
<table_name>;
```

### Note:

- The *LOCAL* Switch specifies that the data we are loading is available in our Local File System. If the *LOCAL* switch is not used, the hive will consider the location as an HDFS path location.
- The *OVERWRITE* switch allows us to overwrite the table data.

Let's make a CSV(Comma Separated Values) file with the name *data.csv* since we have provided ',' as a field terminator while creating a table in the hive. We are creating this file in our local file system at */home/dikshant/Documents* for demonstration purposes.

### Command:

```
cd /home/dikshant/Documents // To change the directory
```

```
touch data.csv // use to create data.csv file
```

```
nano data.csv // nano is a linux command line editor to edit files
```

```
cat data.csv // cat is used to see content of file
```

```
dikshant@dikshant:~$ cd /home/dikshant/Documents/
dikshant@dikshant:~/Documents$ touch data.csv
dikshant@dikshant:~/Documents$ nano data.csv
dikshant@dikshant:~/Documents$ cat data.csv
Ganesh,4,85
Chandan,5,65
Bhavani,6,87
dikshant@dikshant:~/Documents$ █
```

*LOAD DATA* to the student hive table with the help of the below command.

```
LOAD DATA LOCAL INPATH '/home/dikshant/Documents/data.csv' INTO TABLE student;
```

```
hive> LOAD DATA LOCAL INPATH '/home/dikshant/Documents/data.csv' INTO TABLE student;
Loading data to table default.student
OK
Time taken: 2.617 seconds
hive> 
```

Let's see the *student* table content to observe the effect with the help of the below command.  
SELECT \* FROM student;

```
hive> SELECT * FROM student;
OK
Dikshant      1      95.0
Akshat 2      96.0
Dhruv 3       90.0
Ganesh 4      85.0
Chandan 5     65.0
Bhavani 6     87.0
Time taken: 1.24 seconds, Fetched: 6 row(s)
hive> 
```

We can observe that we have successfully added the data to the *student* table.

## EXPERIMENT-5

**To create HDFS tables and load them in Hive and  
Implement joining of tables in Hive.**

### Requirement

You have two table named as A and B. and you want to perform all types of join in hive. It will help you to understand, how join works in hive.



### Solution

#### Step 1: Input Files

Download file A and B from here. And place them into a local directory. File A and B are the comma delimited file, please refer below :-

A	
Id	Type
1	accounts
2	science
3	statistics
4	computer
5	computer

B	
Id	Type
1	accounts
3	science
4	statistics
8	agriculture
9	finance

I am placing these files into local directory 'sample\_files'

```
cd sample_files
```

```
ls -R *
```

Step 2: Loading the files into Hive.

To load the files into hive, Let's first put these files into hdfs location using below commands.

```
hadoop fs -mkdir bdps/sample_files
```

```
hadoop fs -mkdir bdps/sample_files/A
```

```
hadoop fs -mkdir bdps/sample_files/B
```

```
hadoop fs -put A/A.txt bdps/sample_files/A/
```

```
hadoop fs -put B/B.txt bdps/sample_files/B/
```

```

[root@sandbox-hdp sample_data]# ls -R *
A:
A.txt

B:
B.txt

[root@sandbox-hdp sample_data]# hadoop fs -mkdir bdps/sample_files
[root@sandbox-hdp sample_data]# hadoop fs -mkdir bdps/sample_files/A
[root@sandbox-hdp sample_data]# hadoop fs -mkdir bdps/sample_files/B
[root@sandbox-hdp sample_data]# hadoop fs -put A/A.txt bdps/sample_files/A/
[root@sandbox-hdp sample_data]# hadoop fs -put B/B.txt bdps/sample_files/B/
[root@sandbox-hdp sample_data]#

```

you can check the files in hdfs using below command.

```

[root@sandbox-hdp sample_data]# hadoop fs -ls -R hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/
drwxr-xr-x  - root hdfs          0 2018-12-15 03:41 hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/
-rw-r--r--  1 root hdfs        56 2018-12-15 03:41 hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/A/A.txt
drwxr-xr-x  - root hdfs          0 2018-12-15 03:41 hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/B
-rw-r--r--  1 root hdfs        58 2018-12-15 03:41 hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/B/B.txt

```

`hadoop fs -ls -R hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/`

Now let's create two hive table A and B for both the files,using below commands:-

```
CREATE SCHEMA IF NOT EXISTS bdp;  
CREATE EXTERNAL TABLE IF NOT EXISTS  
bdp.A(id INT,type STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION  
'hdfs://sandboxhdp.hortonworks.com:8020/user/root/bdps/sample_files/A';
```

```
CREATE EXTERNAL TABLE IF NOT EXISTS  
bdp.B  
(id INT,type STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE
```



## LOCATION

'hdfs://sandboxhdp.hortonworks.com:8020/user/root/bdps/sample\_files/B';

```
hive> CREATE SCHEMA IF NOT EXISTS bdp;  
OK  
Time taken: 9.265 seconds  
hive> CREATE EXTERNAL TABLE IF NOT EXISTS bdp.A  
  > (id INT,  
  > type STRING)  
  > ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ','  
  > STORED AS TEXTFILE  
  > LOCATION 'hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/A';  
OK  
Time taken: 10.909 seconds  
hive>  
  > CREATE EXTERNAL TABLE IF NOT EXISTS bdp.B  
  > (id INT,  
  > type STRING)  
  > ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ','  
  > STORED AS TEXTFILE  
  > LOCATION 'hdfs://sandbox-hdp.hortonworks.com:8020/user/root/bdps/sample_files/B';  
OK  
Time taken: 1.709 seconds  
hive> █
```

Let's check whether data populated correctly or not using below commands :-

```
select * from bdp.A;
```

```
select * from bdp.B;
```

```
hive> select * from bdp.A;
OK
1      accounts
2      science
3      statistics
4      computer
5      computer
Time taken: 1.929 seconds, Fetched: 5 row(s)
hive> select * from bdp.B;
OK
1      accounts
3      science
4      statistics
8      agriculture
9      finance
Time taken: 0.3 seconds, Fetched: 5 row(s)
```

Let's understand join one by one

#### A. Inner Join:

Sometimes it is required to have only common records out of two datasets. Now we have two table A & B, we are joining based on a key which is id.

So in output, only those records which match id with another dataset will come. Rest will be discarded.

Use below command to perform the inner join.

```
select * from  
(select * from bdp.A)T1  
JOIN  
(select * from bdp.B)T2  
ON T1.id=T2.id;
```

Expected output:

Inner Join			
A		B	
Id	Type	Id	Type
1	accounts	1	accounts
3	statistics	3	science
4	computer	4	statistics

Please refer below screen shot for reference.

```
OK
1      accounts      1      accounts
3      statistics    3      science
4      computer      4      statistics
Time taken: 138.971 seconds, Fetched: 3 row(s)
```

As you can see only records which have the same id such as 1, 3, 4 are present in the output, rest have been discarded.

## B. Left Join

this type of join is performed when we want to look up something from other datasets, the best example would be fetching a phone no of an employee from other datasets based on employee code.

Use below command to perform left join.

```
select * from
(select * from bdp.A)T1
```

## LEFT JOIN

(select \* from bdp.B)T2

ON T1.id=T2.id;

Expected output

Left Join			
A		B	
Id	Type	Id	Type
1	accounts	1	accounts
2	science	NULL	NULL
3	statistics	3	science
4	computer	4	statistics
5	computer	NULL	NULL

Now we have all the records of left table A and matched records of table B.

## C. Right Join

This type of join is performed when we want to get all the data of look-up table with only matching records of left table.

Use below command to perform right join.

```
select * from  
(select * from bdp.A)T1
```

RIGHT JOIN

```
(select * from bdp.B)T2
```

```
ON T1.id=T2.id;
```

Expected output

Right Join			
A		B	
Id	Type	Id	Type
1	accounts	1	accounts
3	statistics	3	science
4	computer	4	statistics
NULL	NULL	8	agriculture
NULL	NULL	9	finance

```
OK  
1      accounts      1      accounts  
3      statistics    3      science  
4      computer      4      statistics  
NULL   NULL         8      agriculture  
NULL   NULL         9      finance  
Time taken: 20.646 seconds, Fetched: 5 row(s)
```

Now we have all the records of right table B and matched records of table A.

D.Full Join

Full Join			
A		B	
Id	Type	Id	Type
1	accounts	1	accounts
2	science	NULL	NULL
3	statistics	3	science
4	computer	4	statistics
5	computer	NULL	NULL
NULL	NULL	8	agriculture
NULL	NULL	9	finance

Now we have all matched and unmatched records in output as shown below.

```

OK
1      accounts      1      accounts
2      science NULL   NULL
3      statistics    3      science
4      computer      4      statistics
5      computer      NULL   NULL
NULL   NULL      8      agriculture
NULL   NULL      9      finance
Time taken: 58.211 seconds, Fetched: 7 row(s)

```

When it is needed to get all the matched and unmatched records out of two datasets, we can use full join. All data from left as well as from right datasets will appear in result set. Nonmatching records will have null values in respective columns.



Use below command to perform full join.

```
select * from  
(select * from bdp.A)T1  
FULL JOIN  
(select * from bdp.B)T2  
ON T1.id=T2.id;
```

Expected output

Now we have all matched and unmatched records in output as shown below.

## Wrapping Up


Joins are important when you have to deal with data which are present in more than a table. In real time we get files from many sources which have a relation between them, so to get meaningful information from these data-sets it needs to perform join to get combined result.

To install, deploy & configure Apache Spark Cluster. To Select the fields from the dataset using Spark SQL.  
To explore Spark shell and read from HDFS



Apache Spark follows a master/slave architecture with two main daemons and a cluster manager.

- Master Daemon — (Master/Driver Process)
- Worker Daemon —(Slave Process)
- Cluster Manager



A spark cluster has a single Master and any number of Slaves/Workers. The driver and the executors run their individual Java processes and users can run them on the same horizontal spark cluster or on separate machines.

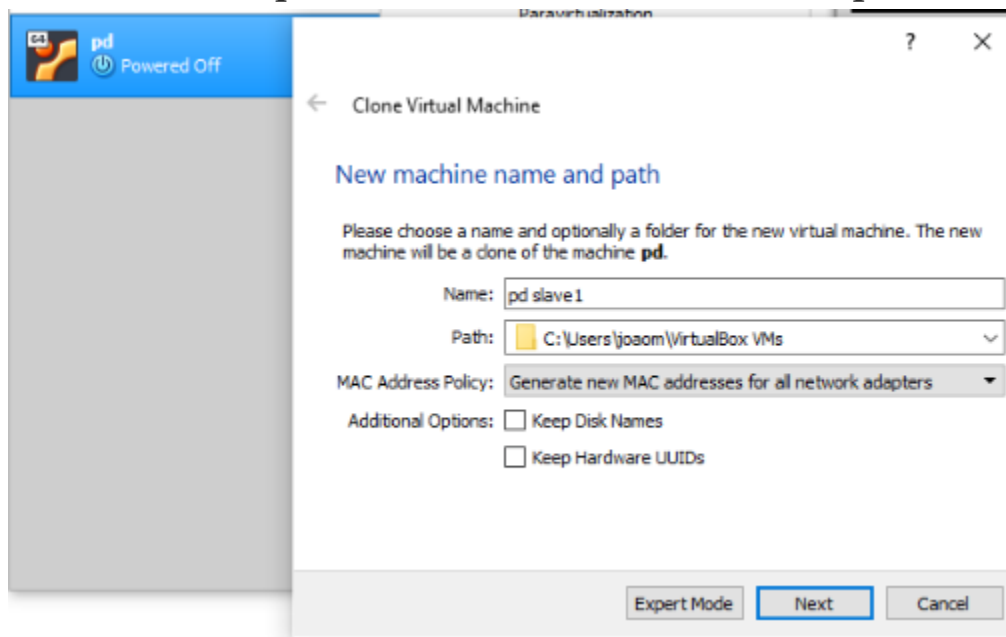
## Pre-requirements

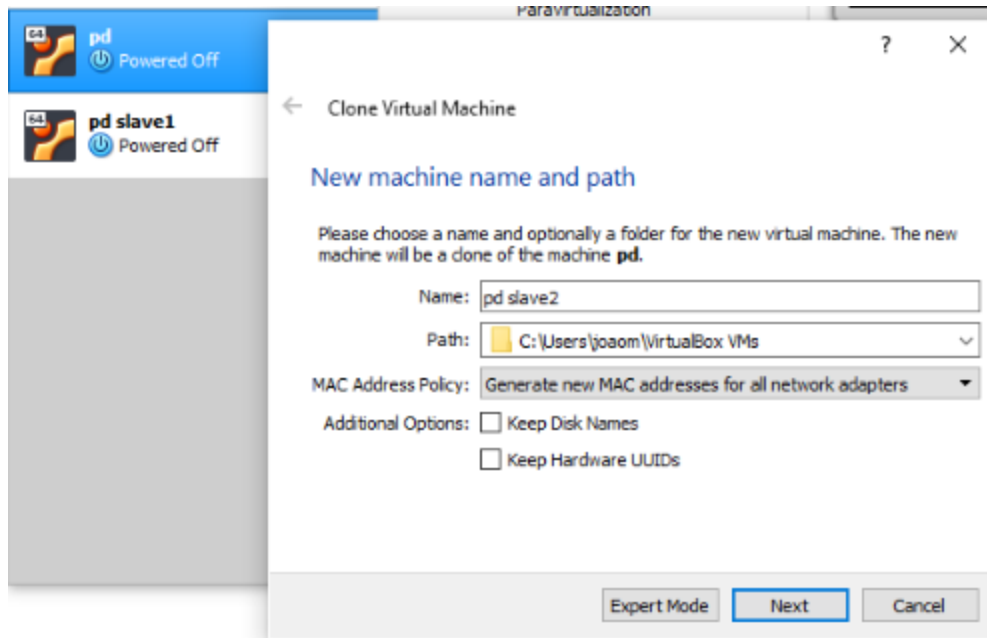
- Ubuntu 18.04 installed on a virtual machine.

## 1st Step:

Create 2 clones of the Virtual Machine you've previously created.

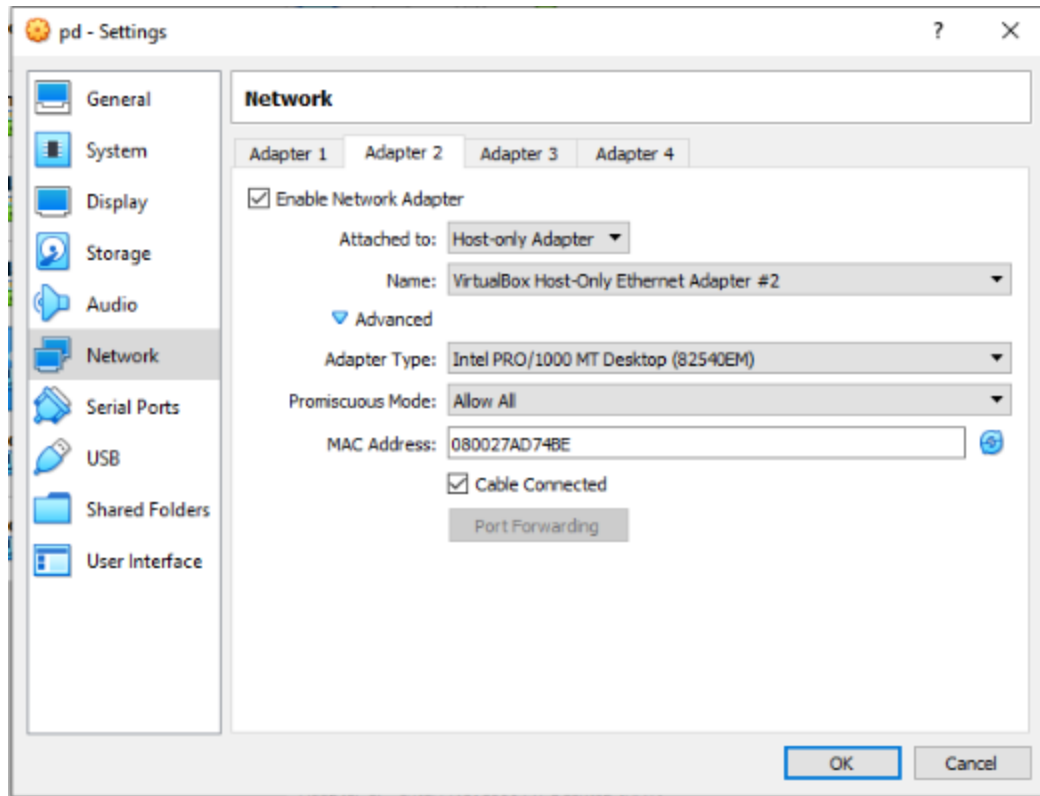
Make sure that you have the option "Generate new MAC addresses for all network adapters selected. Also, choose the option "Full Clone".





## 2nd Step:

Make sure all the VM's have the following network configuration on Apapter 2:



### 3rd Step:

Let's change the hostname on each virtual machine. Open the file and type the name of the machina. Use this command:

```
sudo nano /hostname
```

```
torres@torres-VirtualBox:~$ sudo nano /hostname
[sudo] password for torres:
```

```
torres@torres-VirtualBox: ~  
File Edit View Search Terminal Help  
GNU nano 2.9.3 /hostname Modified  
pd-master  
  
[ New File ]  
^G Get Help    ^O Write Out   ^W Where Is    ^K Cut Text    ^J Justify  
^X Exit        ^R Read File   ^\ Replace     ^U Uncut Text  ^T To Spell
```

```
torres@torres-VirtualBox: ~  
File Edit View Search Terminal Help  
GNU nano 2.9.3 /hostname Modified  
pd-slave1  
  
[ New File ]  
^G Get Help    ^O Write Out   ^W Where Is    ^K Cut Text    ^J Justify  
^X Exit        ^R Read File   ^\ Replace     ^U Uncut Text  ^T To Spell
```



```
torres@torres-VirtualBox: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 /hostname Modified
pd-slave2

[ New File ]
^G Get Help  ^O Write Out  ^W Where Is  ^K Cut Text  ^J Justify
^X Exit      ^R Read File  ^_ Replace   ^U Uncut Text ^T To Spell
```

## 4th Step:

Now let's figure out what our ip address is. To do that just type the command:

```
ip addr
```

```
torres@torres-VirtualBox: ~  
File Edit View Search Terminal Help  
torres@torres-VirtualBox:~$ ip addr  
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group defau  
lt qlen 1000  
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
    inet 127.0.0.1/8 scope host lo  
        valid_lft forever preferred_lft forever  
    inet6 ::1/128 scope host  
        valid_lft forever preferred_lft forever  
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP g  
roup default qlen 1000  
    link/ether 08:00:27:2a:eb:6b brd ff:ff:ff:ff:ff:ff  
    inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic noprefixroute enp0s3  
        valid_lft 86315sec preferred_lft 86315sec  
    inet6 fe80::c66c:d5fa:7179:6628/64 scope link noprefixroute  
        valid_lft forever preferred_lft forever  
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP g  
roup default qlen 1000  
    link/ether 08:00:27:ad:74:be brd ff:ff:ff:ff:ff:ff  
    inet 192.168.205.10/24 brd 192.168.205.255 scope global dynamic noprefixrou  
te enp0s8  
        valid_lft 1116sec preferred_lft 1116sec  
    inet6 fe80::9f4c:576b:e0dc:f4da/64 scope link noprefixroute  
        valid_lft forever preferred_lft forever  
torres@torres-VirtualBox:~$
```

This is on the master VM, as you can see our IP is 192.168.205.10. For you this will be different.

This means that our IP's are:

master: 192.168.205.10

slave1: 192.168.205.11

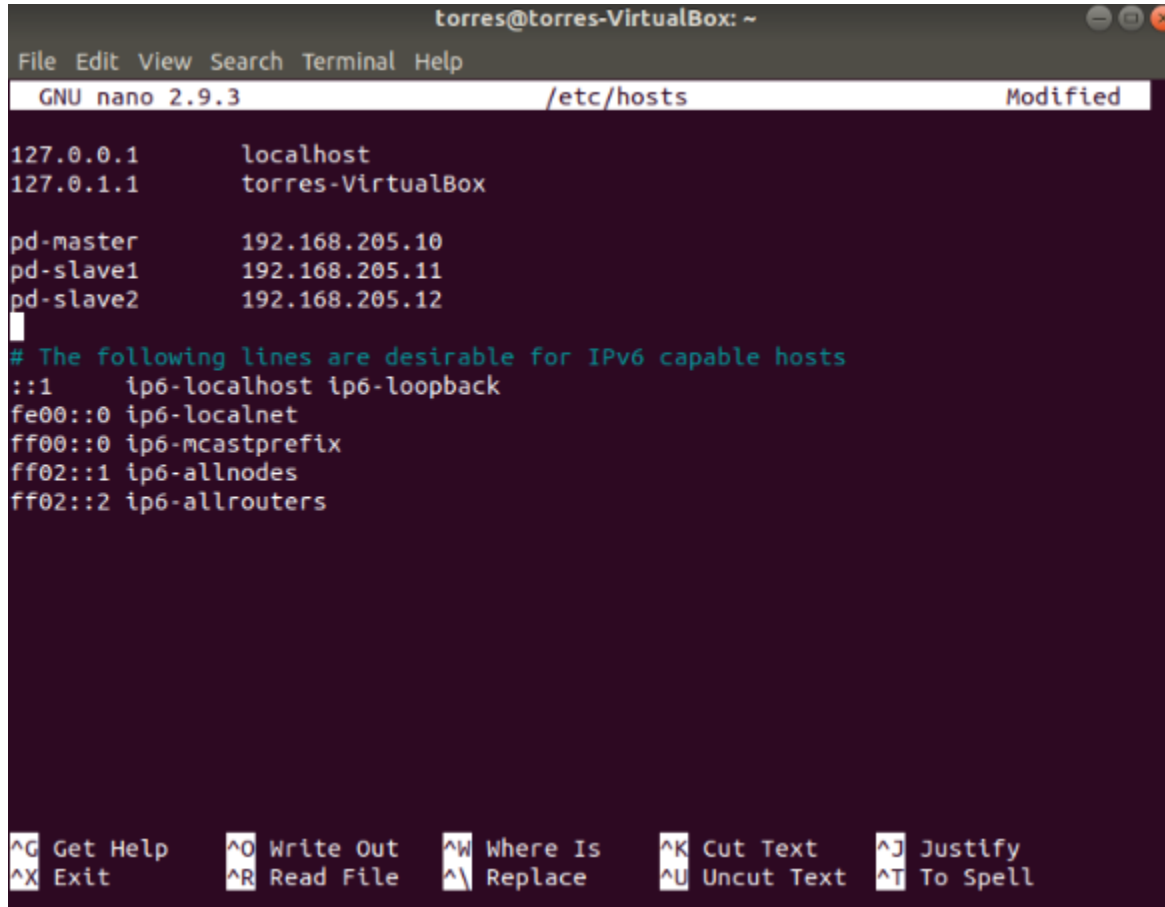
slave2: 192.168.205.12

**5th Step:**

We need to edit the **hosts** file. Use the following command:

```
sudo nano /etc/hosts
```

and add your network information:



```
torres@torres-VirtualBox: ~  
File Edit View Search Terminal Help  
GNU nano 2.9.3 /etc/hosts Modified  
127.0.0.1 localhost  
127.0.1.1 torres-VirtualBox  
  
pd-master 192.168.205.10  
pd-slave1 192.168.205.11  
pd-slave2 192.168.205.12  
|  
# The following lines are desirable for IPv6 capable hosts  
::1 ip6-localhost ip6-loopback  
fe00::0 ip6-localnet  
ff00::0 ip6-mcastprefix  
ff02::1 ip6-allnodes  
ff02::2 ip6-allrouters  
  
^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify  
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell
```

## 6th Step:

In order for the machines to assimilate the previous steps we need to reboot them. Use the following command in all of them:

```
sudo reboot
```

## 7th Step:

Do this step on all the Machines, master and slaves.

Now, in order to install Java we need to do some things. Follow these commands and give permission when needed:

```
$ sudo apt-get install software_properties_common  
$ sudo add-apt-repository ppa:webupd8team/java  
$ sudo apt-get update  
$ sudo apt-get install openjdk-11-jdk
```

To check if java is installed, run the following command.

```
$ java -version
```

```
torres@pd-master:~$ java -version  
openjdk version "11.0.6" 2020-01-14  
OpenJDK Runtime Environment (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1)  
OpenJDK 64-Bit Server VM (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1, mixed mo  
de)  
torres@pd-master:~$
```

## 8th Step:

Now let's install Scala on the master and the slaves. Use this command:

```
$ sudo apt-get install scala
```


```
torres@pd-master:~$ sudo apt-get install scala  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done
```

To check if Scala was correctly installed run this command:

```
$ scala -version
```

```
torres@pd-master:~$ scala -version  
Scala code runner version 2.11.12 -- Copyright 2002-2017, LAMP/EPFL  
torres@pd-master:~$
```

As you can see, Scala version 2.11.12 is now installed on my machine.



# Scala

## 9th Step:

We will configure SSH, but this step in on master only.

We need to install the **Open SSH Server-Client**, use the command:

```
$ sudo apt-get install openssh-server openssh-client
```

```
torres@pd-master:~$ sudo apt-get install openssh-server openssh-client
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  ncurses-term openssh-sftp-server ssh-import-id
Suggested packages:
  keychain libpam-ssh monkeysphere ssh-askpass molly-guard rssh
The following NEW packages will be installed:
  ncurses-term openssh-server openssh-sftp-server ssh-import-id
The following packages will be upgraded:
  openssh-client
1 upgraded, 4 newly installed, 0 to remove and 466 not upgraded.
Need to get 637 kB/1251 kB of archives.
After this operation, 5316 kB of additional disk space will be used.
Do you want to continue? [Y/n]
```

Now generate key pairs:

```
$ ssh-keygen -t rsa -P ""
```

```

torres@pd-master:~$ ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/home/torres/.ssh/id_rsa):
Your identification has been saved in /home/torres/.ssh/id_rsa.
Your public key has been saved in /home/torres/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:OZttQJwxh/PurAbQBxYtns7nhzAAYor9vfB9lNY6uw torres@pd-master
The key's randomart image is:
+----[RSA 2048]-----+
|  .  .o+..  |
|..  .  o+=   |
|o  o...+o    |
| . ..+.. ... |
|o.  o.o S.+  |
|o .  +o. 0.. |
|  . .==o+    |
|  .o+++=o    |
|  .+=oE.     |
+-----[SHA256]-----+
torres@pd-master:~$

```

Use the following command in order to make this key an authorized one:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

```

torres@pd-master:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
torres@pd-master:~$

```

Now we need to copy the content of `.ssh/id_rsa.pub` (of master) to `.ssh/authorized_keys` (of all the slaves as well as master). Use these commands:

```

ssh-copy-id user@pd-master
ssh-copy-id user@pd-slave1
ssh-copy-id user@pd-slave2

```

```

torres@pd-master:~$ ssh-copy-id torres@pd-slave1
The authenticity of host 'pd-slave1 (192.168.205.11)' can't be established.
ECDSA key fingerprint is SHA256:mG56/CnM1qav/9cS7HP8hXNNOEYzzol/7JLKq4cTLbE.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter
out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are promp
ted now it is to install the new keys
torres@pd-slave1's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'torres@pd-slave1'"
and check to make sure that only the key(s) you wanted were added.

torres@pd-master:~$

```

```
torres@pd-master:~$ ssh-copy-id torres@pd-slave2
The authenticity of host 'pd-slave2 (192.168.205.12)' can't be established.
ECDSA key fingerprint is SHA256:vchan4ALBXytJSnv+3L+nChJiQ1GmER3bWpXnMPy4k0.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter
out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are promp
ted now it is to install the new keys
torres@pd-slave2's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'torres@pd-slave2'"
and check to make sure that only the key(s) you wanted were added.

torres@pd-master:~$
```

```
torres@pd-master:~$ ssh-copy-id torres@pd-slave2
The authenticity of host 'pd-slave2 (192.168.205.12)' can't be established.
ECDSA key fingerprint is SHA256:vchan4ALBXytJSnv+3L+nChJiQ1GmER3bWpXnMPy4k0.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter
out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are promp
ted now it is to install the new keys
torres@pd-slave2's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'torres@pd-slave2'"
and check to make sure that only the key(s) you wanted were added.
```

Let's check if everything went well, try to connect to the slaves:

```
$ ssh slave01
$ ssh slave02
```

```
torres@pd-master:~$ ssh pd-slave1
Welcome to Ubuntu 18.04.2 LTS (GNU/Linux 4.18.0-15-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

 * Canonical Livepatch is available for installation.
   - Reduce system reboots and improve kernel security. Activate at:
     https://ubuntu.com/livepatch

463 packages can be updated.
250 updates are security updates.

Your Hardware Enablement Stack (HWE) is supported until April 2023.
torres@pd-slave1:~$
```

```
torres@pd-master:~$ ssh pd-slave2
Welcome to Ubuntu 18.04.2 LTS (GNU/Linux 4.18.0-15-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

 * Canonical Livepatch is available for installation.
   - Reduce system reboots and improve kernel security. Activate at:
     https://ubuntu.com/livepatch

463 packages can be updated.
250 updates are security updates.

Your Hardware Enablement Stack (HWE) is supported until April 2023.
torres@pd-slave2:~$
```

As you can see everything went well, to exit just type the command:

```
exit
```

```
torres@pd-slave2:~$ exit
logout
Connection to pd-slave2 closed.
```

## 10th Step:

Now we Download the latest version of Apache Spark.

**NOTE: Everything inside this step must be done on all the virtual machines.**

Use the following command :

```
$ wget
http://www-us.apache.org/dist/spark/spark-2.4.4/spark-2.4.4-bin-
hadoop2.7.tgz
```

This is the most recent version as of the writing of this article, it might have changed if you try it later. Anyway, I think you'll still be good using this one.



```
torres@pd-master:~$ wget http://www-us.apache.org/dist/spark/spark-2.4.4/spark-2.4.4-bin-hadoop2.7.tgz
--2020-02-03 19:41:41-- http://www-us.apache.org/dist/spark/spark-2.4.4/spark-2.4.4-bin-hadoop2.7.tgz
Resolving www-us.apache.org (www-us.apache.org)... 40.79.78.1
Connecting to www-us.apache.org (www-us.apache.org)|40.79.78.1|:80... connected
.
HTTP request sent, awaiting response... 200 OK
Length: 230091034 (219M) [application/x-gzip]
Saving to: 'spark-2.4.4-bin-hadoop2.7.tgz'

.4.4-bin-hadoop2.7.  0%[          ] 856,50K  231KB/s  eta 16m 45s
```

## Extract the Apache Spark file you just downloaded

Use the following command to extract the Spark tar file:

```
$ tar xvf spark-2.4.4-bin-hadoop2.7.tgz
```

```
torres@pd-master:~$ tar xvf spark-2.4.4-bin-hadoop2.7.tgz
spark-2.4.4-bin-hadoop2.7/
spark-2.4.4-bin-hadoop2.7/R/
spark-2.4.4-bin-hadoop2.7/R/lib/
spark-2.4.4-bin-hadoop2.7/R/lib/sparkr.zlp
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/INDEX
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/html/
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/html/R.css
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/html/00Index.html
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/help/
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/help/aliases.rds
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/help/AnIndex
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/help/SparkR.rdx
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/help/SparkR.rdb
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/help/paths.rds
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/worker/
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/worker/worker.R
spark-2.4.4-bin-hadoop2.7/R/lib/SparkR/worker/daemon.R
```

## Move Apache Spark software files

Use the following command to move the spark software files to respective directory (/usr/local/bin)

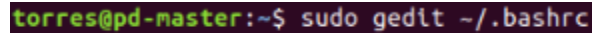
```
$ sudo mv spark-2.4.4-bin-hadoop2.7 /usr/local/spark
```

```
torres@pd-master:~$ sudo mv spark-2.4.4-bin-hadoop2.7 /usr/local/spark
[sudo] password for torres:
torres@pd-master:~$ █
```

## Set up the environment for Apache Spark

Edit the *bashrc* file using this command:

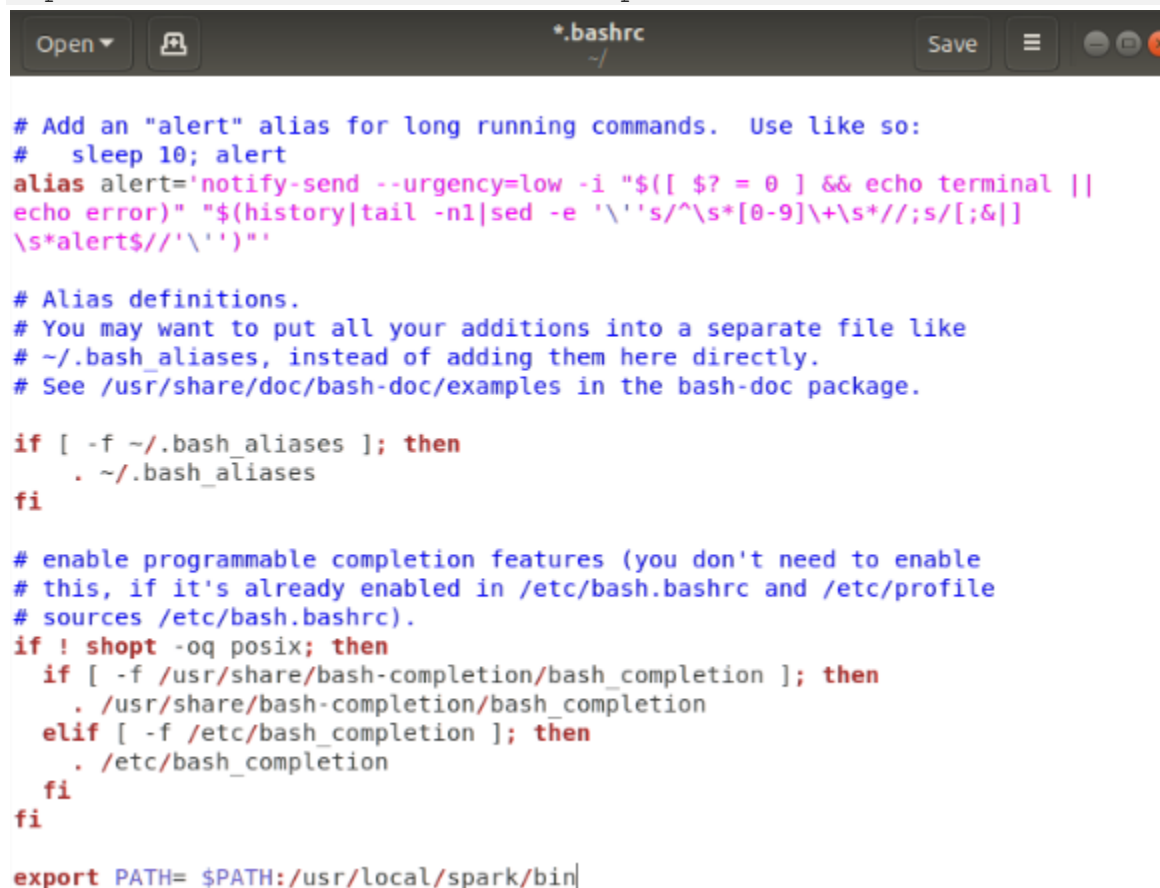
```
$ sudo gedit ~/.bashrc
```



```
torres@pd-master:~$ sudo gedit ~/.bashrc
```

Add the following line to the file. This adds the location where the spark software file are located to the PATH variable.

```
export PATH = $PATH:/usr/local/spark/bin
```



```
Open ▾  *.bashrc  Save  ≡  -  +  x

# Add an "alert" alias for long running commands.  Use like so:
#   sleep 10; alert
alias alert='notify-send --urgency=low -i "${[ $? = 0 ] && echo terminal ||
echo error}" "$(history|tail -n1|sed -e '\''s/^\s*[0-9]\+\s*//;s/[\;&|]
\s*alert$//'\'' )"'

# Alias definitions.
# You may want to put all your additions into a separate file like
# ~/.bash_aliases, instead of adding them here directly.
# See /usr/share/doc/bash-doc/examples in the bash-doc package.

if [ -f ~/.bash_aliases ]; then
    . ~/.bash_aliases
fi

# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
if ! shopt -oq posix; then
    if [ -f /usr/share/bash-completion/bash_completion ]; then
        . /usr/share/bash-completion/bash_completion
    elif [ -f /etc/bash_completion ]; then
        . /etc/bash_completion
    fi
fi

export PATH= $PATH:/usr/local/spark/bin|
```

**Note:** this screenshot has a mistake, when you're doing this don't leave a space like I did. Just write "PATH=\$PATH".

Now we need to use the following command for sourcing the *~/.bashrc* file:

```
$ source ~/.bashrc
```

```
torres@pd-master:~$ source ~/.bashrc
torres@pd-master:~$
```

## 11th Step:

Apache Spark Master Configuration (do this step on the Master VM only)

### Edit spark-env.sh

Move to spark *conf* folder and create a copy of the template of *spark-env.sh* and rename it.

```
$ cd /usr/local/spark/conf
$ cp spark-env.sh.template spark-env.sh
torres@pd-master:/usr/local/spark/conf$ cp spark-env.sh.template spark-env.sh
torres@pd-master:/usr/local/spark/conf$
```

Now edit the configuration file *spark-env.sh*.

```
$ sudo vim spark-env.sh
```

And add the following parameters:

```
export SPARK_MASTER_HOST='<MASTER-IP>'export
JAVA_HOME=<Path_of_JAVA_installation>
```

### Add Workers

Edit the configuration file *slaves* in (*/usr/local/spark/conf*).

```
$ sudo nano slaves
torres@pd-master:/usr/local/spark/conf$ sudo nano slaves
```

And add the following entries.

```
pd-master
pd-slave01
pd-slave02
```

```
torres@pd-master: /usr/local/spark/conf
File Edit View Search Terminal Help
GNU nano 2.9.3 slaves Modified
pd-master
pd-slave1
pd-slave2

^G Get Help      ^O Write Out    ^W Where Is     ^K Cut Text     ^J Justify
^X Exit          ^R Read File    ^\ Replace      ^U Uncut Text   ^T To Spell
```

## 12th Step:

**Let's try to start our Apache Spark Cluster, hopefully everything is ok!**

To start the spark cluster, run the following command on master.:

```
$ cd /usr/local/spark
$ ./sbin/start-all.sh
```

```
torres@pd-master:/usr/local/spark/conf$ cd /usr/local/spark
torres@pd-master:/usr/local/spark$
```

```
torres@pd-master:/usr/local/spark$ ./sbin/start-all.sh
starting org.apache.spark.deploy.master.Master, logging to /usr/local/spark/log
s/spark-torres-org.apache.spark.deploy.master.Master-1-pd-master.out
pd-slave1: starting org.apache.spark.deploy.worker.Worker, logging to /usr/loca
l/spark/logs/spark-torres-org.apache.spark.deploy.worker.Worker-1-pd-slave1.out
pd-master: starting org.apache.spark.deploy.worker.Worker, logging to /usr/loca
l/spark/logs/spark-torres-org.apache.spark.deploy.worker.Worker-1-pd-master.out
```

I won't stop it, but in case you want to stop the cluster, this is the command:

```
$ ./sbin/stop-all.sh
```

### 13th Step:

To check if the services started we use the command:

```
$ jps
```

```
torres@pd-master:/usr/local/spark$ jps
9280 Jps
9065 Master
9212 Worker
```

### 14th Step:

Browse the Spark UI to know about your cluster. To do this, go to your browser and type:

```
http://<MASTER-IP>:8080/
```

Spark Master at spark://192.168.205.10:7077 - Mozilla Firefox

Spark Master at spark:// x Firefox Privacy Notice — x +

192.168.205.10:8080

**Spark** 2.4.4 **Spark Master at spark://192.168.205.10:7077**

URL: spark://192.168.205.10:7077  
Alive Workers: 2  
Cores in use: 2 Total, 0 Used  
Memory in use: 2.0 GB Total, 0.0 B Used  
Applications: 0 Running, 0 Completed  
Drivers: 0 Running, 0 Completed  
Status: ALIVE

▼ Workers (2)

Worker Id	Address	State	Cores	Memory
<a href="#">worker-20200203205212-192.168.205.11-32793</a>	192.168.205.11:32793	ALIVE	1 (0 Used)	1024.0 MB (0.0 B Used)
<a href="#">worker-20200203205213-192.168.205.10-44377</a>	192.168.205.10:44377	ALIVE	1 (0 Used)	1024.0 MB (0.0 B Used)

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

As you can see we have 2 Alive Workers, our slaves, which means it's all done!

## Assignment 8

**Title:** To Perform Graph analytics and visualization using Tableau.

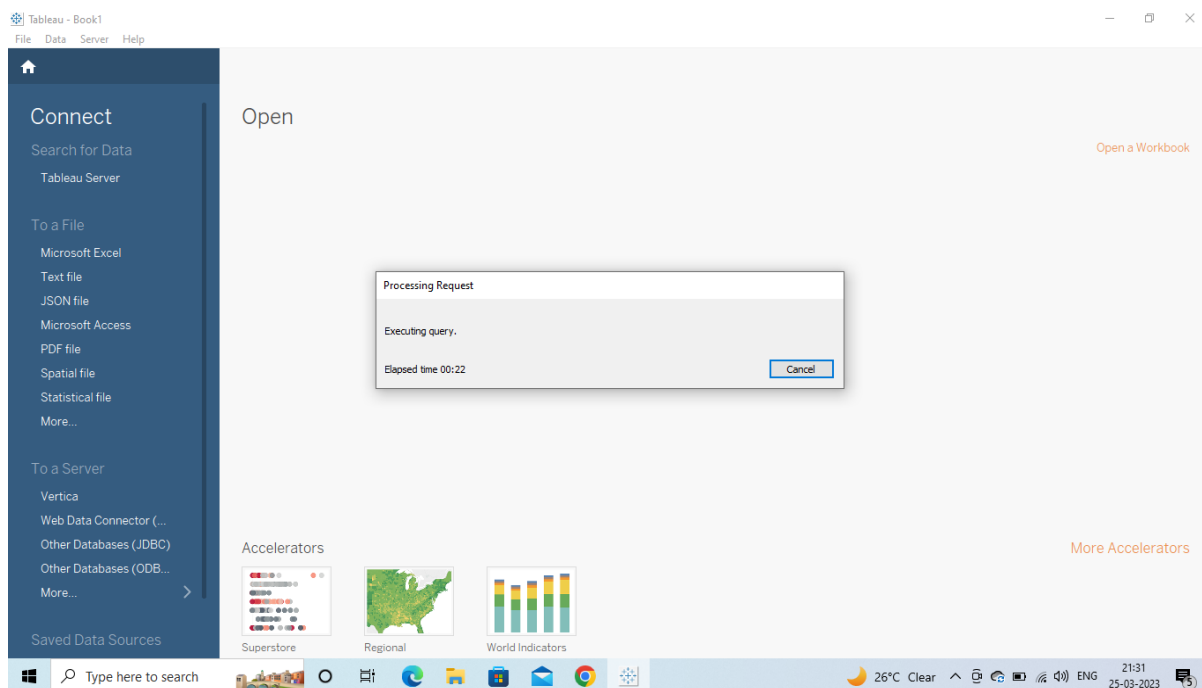
### SOFTWARE REQUIREMENTS:

1. windows computer or Mac computer
2. Tableau

### THEORY:

Tableau is a data analytics and visualization tool used widely in the industry today. Many businesses even consider it indispensable for data-science-related work. Tableau's ease of use comes from the fact that it has a drag and drop interface. This feature helps to perform tasks like sorting, comparing and analyzing, very easily and fast. Tableau is also compatible with multiple sources, including Excel, SQL Server, and cloud-based data repositories which makes it an excellent choice for Data Scientists.

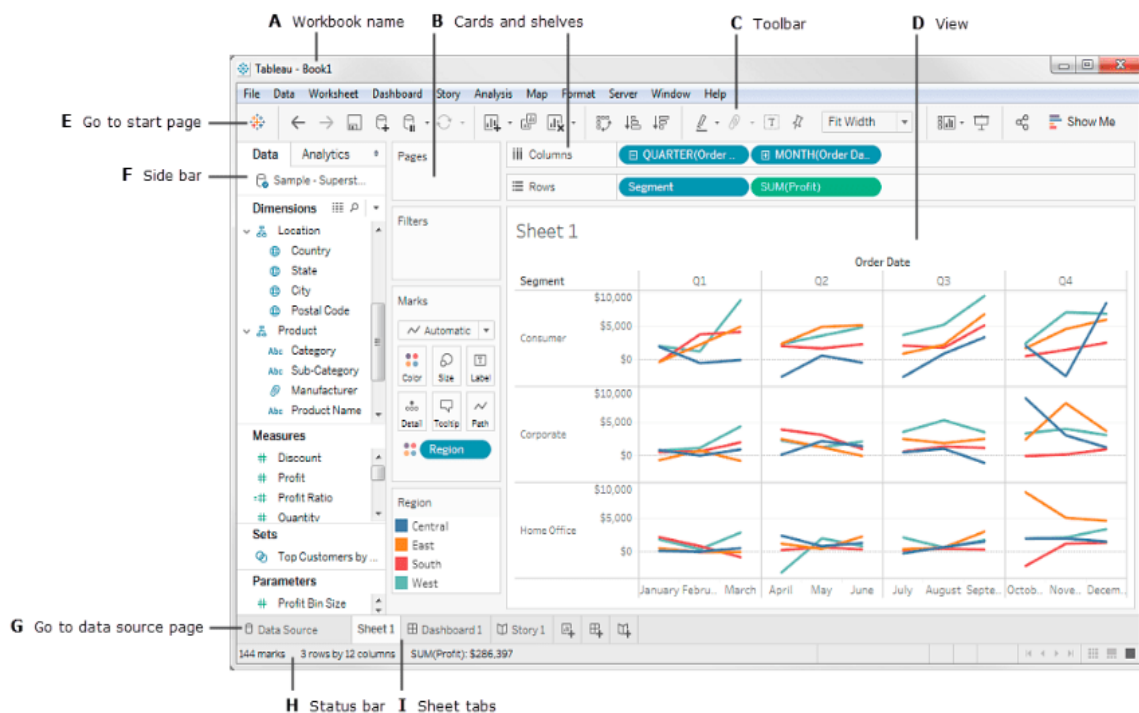
**Step1:** Download and install the Tableau.



**image 1: Tableau**

## Tableau Workspace

The Tableau workspace is a collection of worksheets, menu bar, toolbar, marks card, shelves and a lot of other elements. Sheets can be worksheets, dashboards, or stories. The image below highlights the major components of the workspace. However, more familiarity will be achieved once we work with actual data.



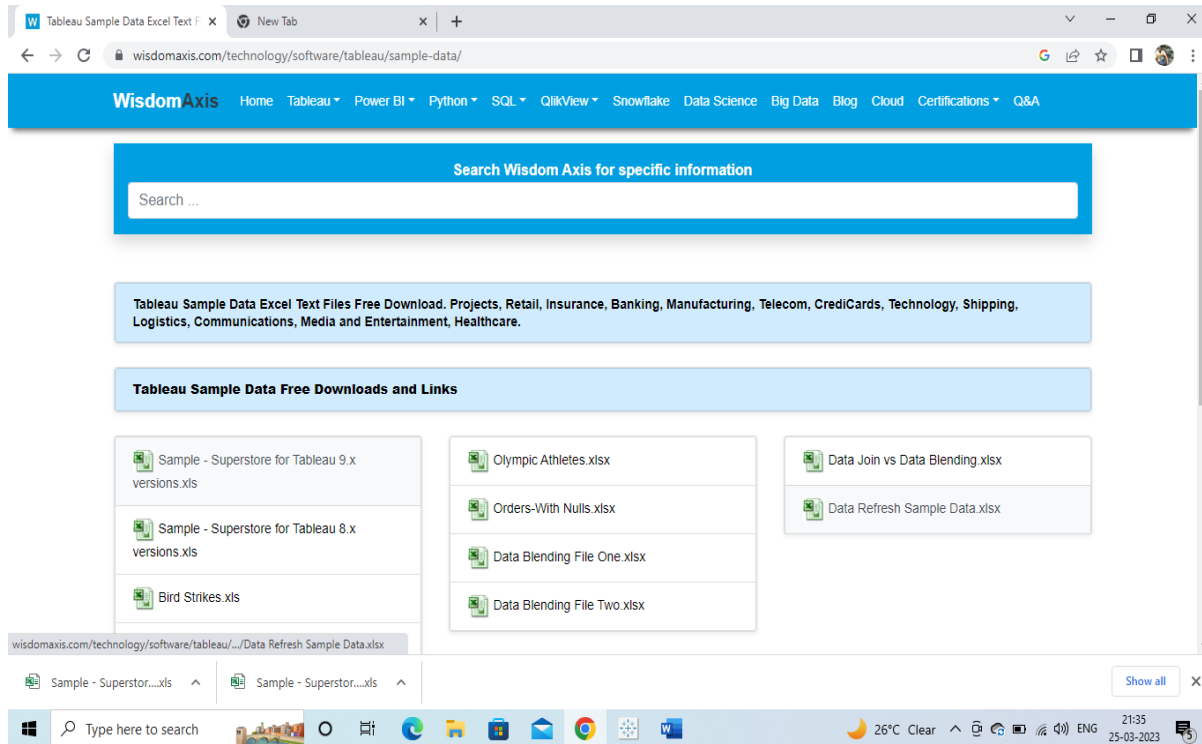
**Image 2: Tablue Workspace**

**Step2:** Download the sample data file in excel.


Tableau is compatible with a lot of data sources. The data sources supported by Tableau appear on the left side of the opening screen. Some




commonly used data sources are excel, text file, relational database or even on a server.



**Image 3: Sample Data for Tableau**

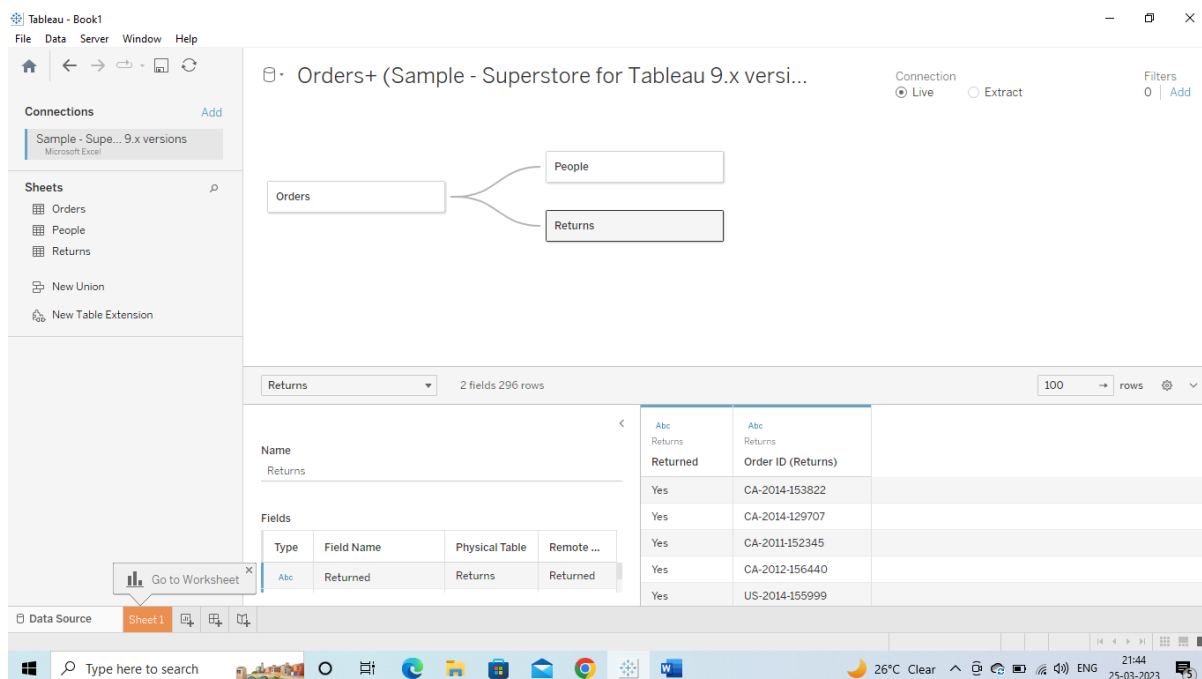
**Step 3:** Select open menu  open downloaded data file



**Image 4: Selection Of sample data file**

**Step 4:** Create the connection between Data fields from the database.

Form sample data drag and drop the files. Here people, return and order select and drop on the workspace. Create the connection between them.



## Image 5: Create Data field connection

Go to sheet1 for visualization of the data. Different visualization available in form of graph.

**Step 5:** Select row and column as customer and their respective country. Field selection done using available data fields.

**Step 6:** Select chart type for more visualization of the data.

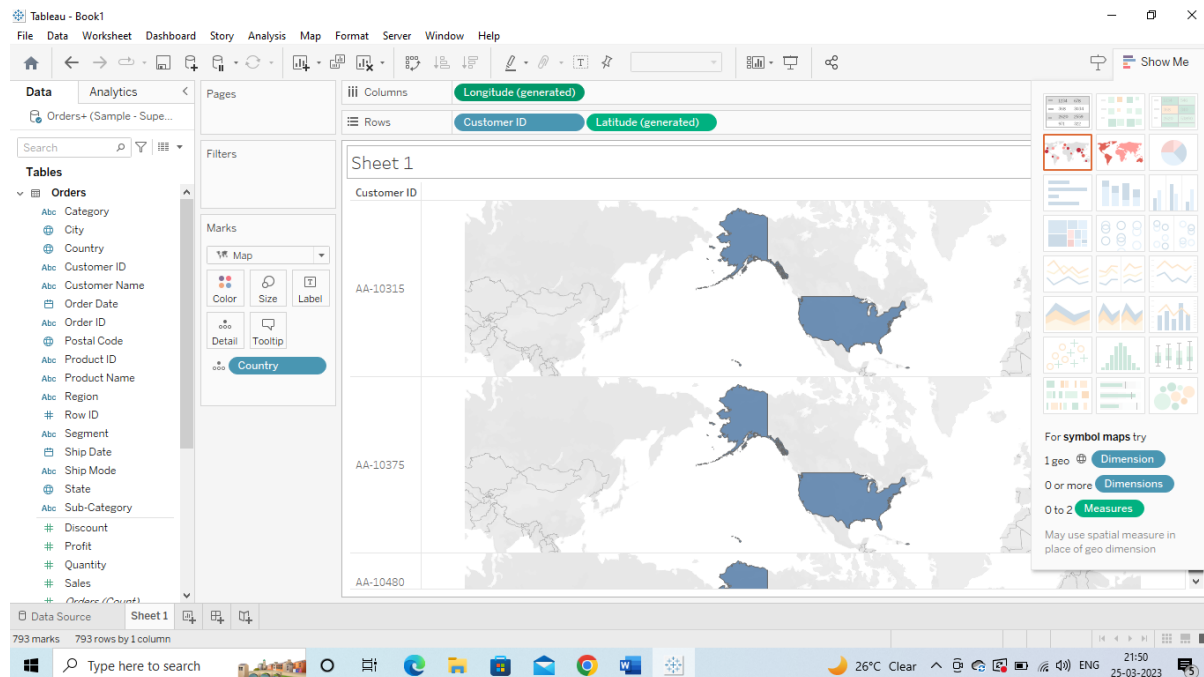
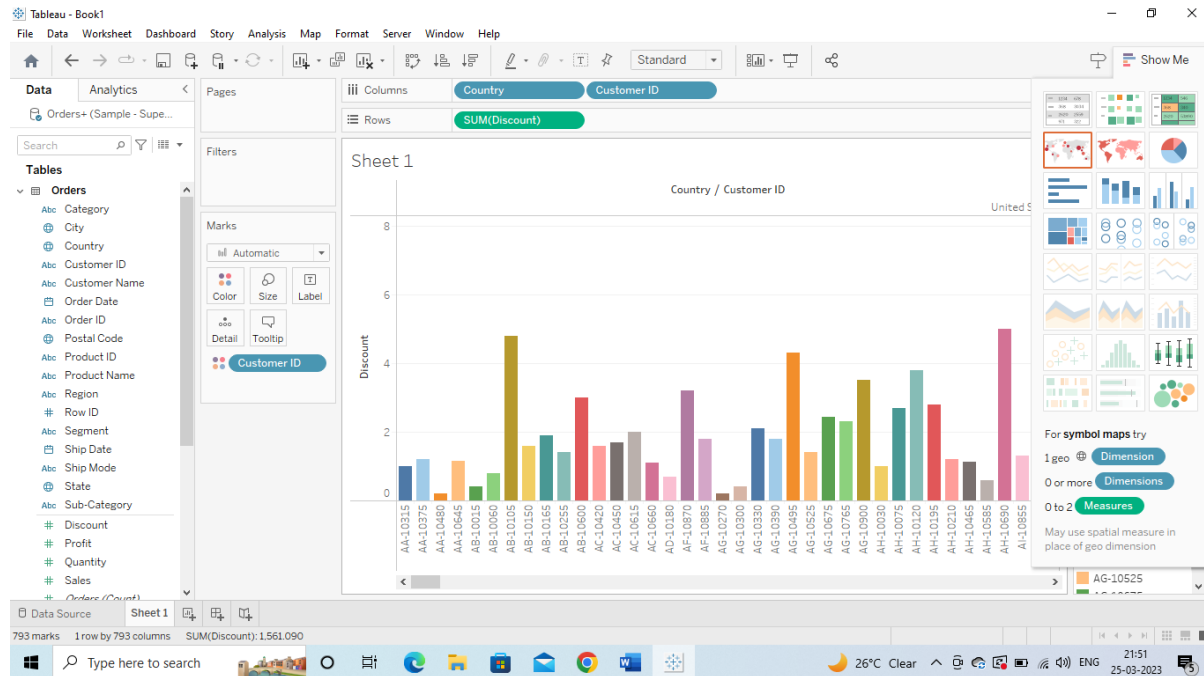


Image 6: Data visualization using map



**Image 7: Data visualization using Bar Chart**

**Conclusion:** In such way Tableau Visualization Though the dataset is complex or the dataset is very big, in tableau, we can create dashboards very easily and within less time.